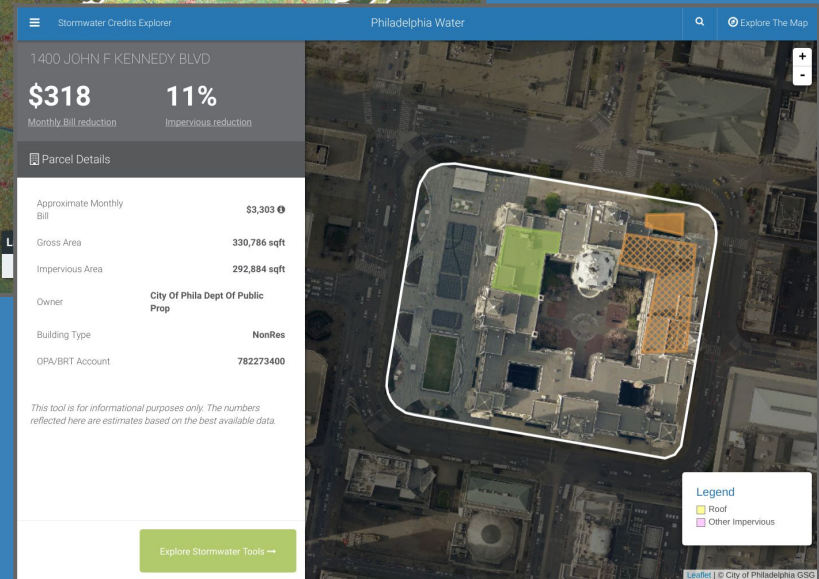
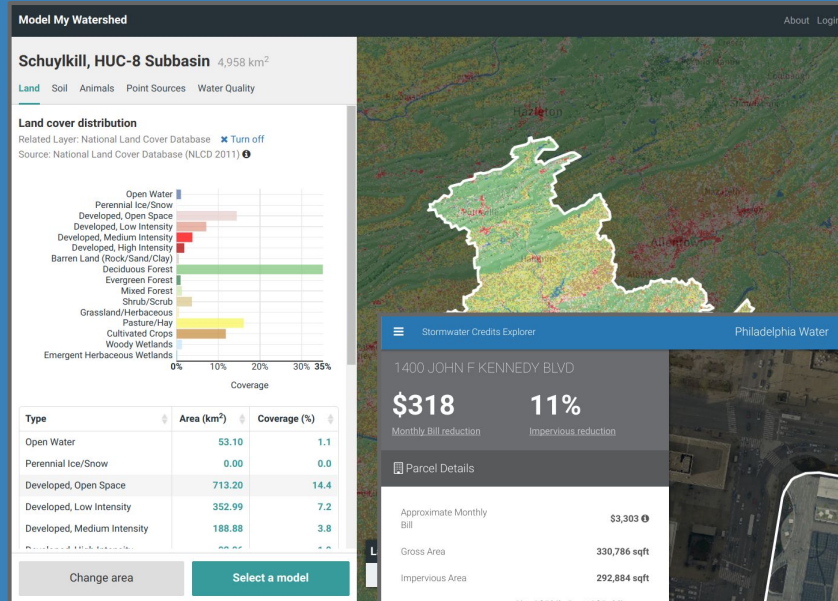


Raster processing on Serverless Architecture

Using python tools on AWS Lambda

Hello.



What is Serverless Architecture?

What is Serverless Architecture?

- **Stateless function as unit of work**
- **Event driven**
- **Cost per execution**
- **No provisioning, limited management**
- **Auto-scaling, fault tolerant by default**

AWS Lambda



AWS Lambda - how it works

- **Write a function that handles events**
- **Implement workflow as usual**
- **Zip code and dependencies**
- **Upload and register/update**

AWS Lambda - what you get

- **A runtime (python, node, java, .NET) on linux**
- **Compute resources (bundled RAM/CPU)**
- **Event integrations**
- **Web UI, CLI , 3rd party tools**
- **Lots of limitations**

AWS Lambda - execution limits

- **Up to 1.5GB memory (proportional CPU)**
- **512MB / tmp space**
- **300s timeout**
- **6MB limit response/request**
- **Cap on threads/process/file descriptors**

AWS Lambda - deployment

- **Code + Dependencies bundled**
 - **50mb compressed**
 - **250mb uncompressed**
- **Build dependencies which target AMZ Linux**
- **Deployment size impacts cold start time**

AWS Lambda - costs

- **Charged per request + per 100ms**
 - **\$0.20 / 1MM requests**
 - **\$0.00001667 / GB-s**
- **Perpetual free-tier (per month)**
 - **First million requests**
 - **First 400,000 GB-s**

When to use

- **Latency tolerant**
- **Stateless**
- **Traffic is bursty**
- **Relatively short-running**
- **Modest resource requirements**

Example Use Cases

- **ETL**
- **“Micro-services”**
- **Raster statistics**
- **Raster tiling**

Python Raster Tools

rasterio & numpy (and friends)

Rasterio

- **Idiomatic python API over GDAL**
- **Reads raster data into numpy arrays**
- **Raster processing utilities**
- **Windowed reads**
- **HTTP/S3 reads**

Numpy

- **Large, multidimensional arrays (aka, rasters)**
- **Not spatial**
- **Typically fast**
- **Extensive community and ecosystem**

And more

- **Shapely: vector operations**
- **Pyproj: reprojection**
- **Pillow: image generation**


How it's done

Optimize your rasters

- **Target S3**
- **Compress**
- **Internal tiles and and overviews**
- **Use VRTs (with /vsi3)**

Optimize your reads

- **Windowed reads**




becomes...

2	2	2	6
-	12	6	6
-	12	12	12
-	-	-	-

Optimize your reads


- **Windowed reads**
- **Chunk windows**




Optimize your reads

- **Windowed reads**
- **Chunk windows**
- **Parallelize**

Input



Chunks



Optimize your reads

- **Windowed reads**
- **Chunk windows**
- **Parallelize**
- **Decimated reads + overviews for tiles**

Serverless Framework

- **NodeJS based CLI**
- **YAML config file to represent resources**
- **Don't use it to package your deps**

Dependencies

- **Use a virtualenv**
- **Build in an Amzn Linux Container (if needed)**
- **Create custom zip**
 - **Symlink duplicated C libraries**
 - **High Compression**

Demos!

<http://raster.surge.sh/>

Summary

- **Use Lambda for geoprocessing**
- **Also, don't use it**
- **Don't abandon good engineering practices**

Thanks!

Matt McFarland
mmcfarland@azavea.com



<https://github.com/mmcfarland/foss4g-lambda-demo>